

Industry Daily Report — Source Checklist & Diagnosis

Against the Project Brief of 28 Jul 2026 (66 embedded links) · live pipeline state as of 26 August 2026

66
links in the client brief

73
sources registered (68 active)

49
working — delivering articles

16
working but quiet this week

3
failing (visible, tracked)

5
off — walled or pending

108
articles captured, last 24 h

853
articles captured, last 7 days

05:00
report send (AEST), reads the finished prior day

3×/day
quality passes: 06:00 · 12:00 · 23:00

Every link in the brief maps to a registered source below — 66 of 66. The registry carries seven more than the brief because several brief links share a publisher (split into sections) and the newsletter intake adds sources of its own. A source marked QUIET fetched successfully and found the publisher had released nothing; that state is recorded and watched, never hidden behind an "ok".

Australian News

SOURCE	METHOD	STATUS	ARTICLES 7D/24H	LAST ARTICLE	ERROR	REASON / REMEDY
ABC News — Climate Change abc.net.au	scrape	WORKING	8 / 0	2026-08-24		Fetches and parses cleanly on schedule.
ABC News — Energy abc.net.au	scrape	WORKING	7 / 0	2026-08-25		Fetches and parses cleanly on schedule.
Australian Financial Review — Energy afr.com	scrape	WORKING	38 / 6	2026-08-25		Listing scrapes cleanly; article bodies are paywalled, so summaries come from the listing. Most items carry no machine-readable date - the article-page date reader now fills them in.
Newsletter — unrecognised sender	email	OFF — pending	0 / 0	—		Catch-all for newsletters from senders not yet mapped to a source. Fills only when such mail arrives.
The Australian — Mining & Energy theaustralian.com.au	email	OFF — account	0 / 0	—	login required	Subscription site. Newsletter subscribed to the intake mailbox; articles arrive by email (tier 0), not by scraping.
The Conversation AU — Environment & Energy theconversation.com	scrape	WORKING	15 / 2	2026-08-26		Healthy scrape; some articles carry no meta date, filled by the article-page reader.

Industry Updates

SOURCE	METHOD	STATUS	ARTICLES 7D/24H	LAST ARTICLE	ERROR	REASON / REMEDY
AER — Registers Status aer.gov.au	scrape	FAILING	0 / 0	—	empty x494	Same class as vic-premier: client-rendered register, no feed, no links in the HTML. Unsolved and visibly failing (494+ consecutive).
Victorian Premier — Energy & Climate Media premier.vic.gov.au	scrape	FAILING	0 / 0	2026-08-01	empty x428	Articles load by in-page XHR; the listing HTML carries no links for any parser. Unsolved - needs its JSON endpoint captured. Fails visibly (428+ consecutive) rather than pretending health.
ACCC — Energy acc.gov.au	scrape	WORKING	5 / 0	2026-08-23		Fetches and parses cleanly on schedule.
AEMC — News Centre aemc.gov.au	rss	WORKING	3 / 0	2026-08-21		Fetches and parses cleanly on schedule.
AEMC — Rule Changes / Registers aemc.gov.au	rss	QUIET	0 / 0	2026-08-05		Rule-change register; publishes in bursts around determinations. Quiet weeks are normal.

AEMO aemo.com.au	scrape	WORKING	52 / 0	2026-08-24	Fetches heavily but the newsroom links site navigation alongside stories - a podcast page, a scholarship, market explainers. The news-vs-reference judge now holds those out of the sheet.
AER — News aer.gov.au	scrape	WORKING	53 / 0	2026-08-23	Fetches and parses cleanly on schedule.
ARENA — Knowledge Bank arena.gov.au	rss	QUIET	0 / 0	2026-08-11	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
ARENA — News arena.gov.au	rss	QUIET	0 / 0	2026-08-04	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
Australian Alliance for Energy Productivity (AIDC) aidc.org.au	scrape	QUIET	0 / 0	2026-07-31	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
Australian Building Codes Board abcb.gov.au	scrape	QUIET	0 / 0	2026-08-06	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
Australian Energy Council energycouncil.com.au	rss	WORKING	3 / 0	2026-08-25	Fetches and parses cleanly on schedule.
Australian Hydrogen Council aidc.org.au	scrape	QUIET	0 / 0	—	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
CEFC — Insights cefc.com.au	scrape	QUIET	0 / 0	2026-08-13	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
CEFC — Media cefc.com.au	scrape	WORKING	4 / 0	2026-08-25	Fetches and parses cleanly on schedule.
Clean Energy Council cleanenergycouncil.org.au	scrape	WORKING	40 / 0	2026-08-23	Every story was once lost to a "Find out more" link text (scored as too short); fixed via slug-derived titles. Publishes no date in its markup; member-portal login (tier 3) is available for the Monday Megawatt.
Clean Energy Regulator — Media Releases cer.gov.au	scrape	QUIET	0 / 0	2026-07-31	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
Clean Energy Regulator — News cer.gov.au	scrape	WORKING	2 / 1	2026-08-25	Fetches and parses cleanly on schedule.
Clean Energy Regulator — NGER cer.gov.au	scrape	QUIET	0 / 0	—	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
Climate Change Authority climatechangeauthority.gov.au	rss	WORKING	1 / 0	2026-08-20	Fetches and parses cleanly on schedule.
Climate Works Centre climateworkscentre.org	scrape	WORKING	1 / 1	2026-08-25	Fetches and parses cleanly on schedule.
Commercial Building Disclosure cbd.gov.au	rss	QUIET	0 / 0	2026-08-17	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
CSIRO — Energy News csiro.au	scrape	WORKING	3 / 0	2026-08-23	Fetches and parses cleanly on schedule.
DCCEEW dcceew.gov.au	scrape	WORKING	28 / 1	2026-08-25	Every article was once dropped because /about/news/ paths were mistaken for chrome; fixed. 28 articles / 7d since.
DEECA — Energy News energy.vic.gov.au	scrape	QUIET	0 / 0	—	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
DISR industry.gov.au	scrape	WORKING	52 / 3	2026-08-25	Fetches and parses cleanly on schedule.
Eco Generation ecogeneration.com.au	rss	WORKING	11 / 3	2026-08-26	TLS-fingerprint 403 until the Scrapling fetcher; now healthy on RSS (11 / 7d). Flagged by the client on 21 Aug - the miss predates the fix.
Energy Efficiency Council eec.org.au	rss	WORKING	1 / 1	2026-08-25	Fetches and parses cleanly on schedule.
Energy Magazine energymagazine.com.au	rss	WORKING	7 / 0	2026-08-25	Fetches and parses cleanly on schedule.

Energy Rating / GEMS energyrating.gov.au	scrape	WORKING	44 / 0	2026-08-23	Parser found nothing for 71 consecutive fetches until the markup fix; 44 articles / 7d since. Captured pages are often program/registry pages - judged before the sheet.
Energy Users Association of Australia euaa.com.au	rss	WORKING	2 / 1	2026-08-25	Fetches and parses cleanly on schedule.
Energy.gov.au — News energy.gov.au	rss	WORKING	5 / 0	2026-08-24	Fetches and parses cleanly on schedule.
Essential Services Commission — Media esc.vic.gov.au	rss	WORKING	1 / 0	2026-08-25	Fetches and parses cleanly on schedule.
Facility Management Association fma.com.au	scrape	WORKING	1 / 0	2026-08-24	Fetches and parses cleanly on schedule.
Green Building Council of Australia new.gbca.org.au	scrape	WORKING	4 / 1	2026-08-25	Fetches and parses cleanly on schedule.
Infrastructure Magazine — Facilities Management infrastructuremagazine.com.au	rss	WORKING	18 / 5	2026-08-26	Fetches and parses cleanly on schedule.
Investor Group on Climate Change igcc.org.au	rss	WORKING	1 / 0	2026-08-20	Fetches and parses cleanly on schedule.
Minister Bowen — Media Releases minister.dccew.gov.au	scrape	WORKING	16 / 0	2026-08-25	Fetches and parses cleanly on schedule.
NABERS nabers.gov.au	scrape	WORKING	34 / 0	2026-08-25	TLS-fingerprint 403 solved by Scrapling. Publishes no dates; several captured pages are program pages, which the judge now separates from news.
Parliament of Australia — Bills & Legislation aph.gov.au	scrape	WORKING	8 / 0	2026-08-20	Fetches and parses cleanly on schedule.
Property Council of Australia propertycouncil.com.au	rss	WORKING	12 / 5	2026-08-25	Parser silent for 527 consecutive fetches until fixed; now 12 / 7d on RSS.
PV Magazine Australia pv-magazine-australia.com	rss	WORKING	24 / 5	2026-08-26	Flagged by the client on 21 Aug; now RSS plus newsletter (24 / 7d).
RACE for 2030 racefor2030.com.au	rss	QUIET	0 / 0	—	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
RenewEconomy reneweconomy.com.au	rss	WORKING	51 / 12	2026-08-26	Flagged by the client on 21 Aug; now RSS plus the daily newsletter into the intake mailbox (51 / 7d).
RenewEconomy — David Leitch reneweconomy.com.au	rss	QUIET	0 / 0	2026-08-11	Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
SEC Victoria secvictoria.com.au	scrape	WORKING	9 / 0	2026-08-23	Fetches and parses cleanly on schedule.
Smart Energy Council smartenergy.org.au	rss	WORKING	7 / 1	2026-08-26	Fetches and parses cleanly on schedule.
Solar Victoria solar.vic.gov.au	scrape	QUIET	0 / 0	2026-07-31	Quiet since 31 Jul. Publishes rarely; worth a manual glance if it stays silent past mid-September.
The Fifth Estate thefifthestate.com.au	rss	WORKING	9 / 0	2026-08-24	Flagged by the client on 21 Aug; now healthy on RSS (9 / 7d).
VicGrid vicgrid.com.au	scrape	WORKING	1 / 0	2026-08-23	Fetches and parses cleanly on schedule.
Victorian Energy Upgrades — News energy.vic.gov.au	scrape	WORKING	8 / 0	2026-08-23	Fetches and parses cleanly on schedule.

Victorian Energy Upgrades — Program Updates esc.vic.gov.au	rss	QUIET	0 / 0	2026-08-03	Feed publishes rarely; last item 3 Aug. Quiet, not broken.
--	-----	-------	-------	------------	--

International Industry Updates

SOURCE	METHOD	STATUS	ARTICLES 7D/24H	LAST ARTICLE	ERROR	REASON / REMEDY
S&P Global — Natural Gas spglobal.com	sitemap	FAILING	0 / 0	—	empty x12	Sitemap fetches (1.1 MB) but yields no recent articles. Under watch: either the section is quiet or the sitemap is the wrong one - not yet proven either way.
Energy Tracker Asia energytracker.asia	rss	QUIET	0 / 0	2026-08-03		Fetch succeeds; the publisher has simply released nothing in the window. Quiet is watched, not hidden.
European Commission — Energy energy.ec.europa.eu	scrape	WORKING	1 / 0	2026-08-19		Fetches and parses cleanly on schedule.
IEA iea.org	scrape	OFF — account	0 / 0	—	login required	Account wall (HTTP 403 to every fetcher). Replaced by newsletter intake; the old scrape source is retired at 527 consecutive failures.
IEA — Energy Efficiency iea.org	scrape	OFF — pending	0 / 0	—		Successor to the IEA scrape, awaiting the newsletter route. Never run yet.
IEEFA ieefa.org	rss	WORKING	14 / 4	2026-08-25		Fetches and parses cleanly on schedule.
IRENA irena.org	scrape	WORKING	1 / 0	2026-08-19		Fetches and parses cleanly on schedule.
NZ Herald — Energy nzherald.co.nz	scrape	WORKING	7 / 1	2026-08-25		Fetches and parses cleanly on schedule.
Renewables Now renewablesnow.com	scrape	WORKING	167 / 33	2026-08-25		Highest-volume source (167 articles / 7d). Rate-limits (HTTP 429) under parallel fetching - per-host politeness is mandatory. No date in its markup; dates are read from each article page.
Reuters — Carbon Markets reuters.com	scrape	OFF — account	0 / 0	—	login required	Account wall (HTTP 401). A genuine login, not a fingerprint block. Route is the free Reuters newsletter into the intake mailbox; inactive until that lands so it cannot report false health.
S&P Global — Crude Oil spglobal.com	sitemap	WORKING	7 / 7	2026-08-26		Same block as the parent section; read from its own sitemap. Working since 26 Aug (7 articles, including the exact stories the client flagged missing).
S&P Global — Energy Transition spglobal.com	scrape	WORKING	44 / 0	2026-08-24		Older scrape source. TLS-fingerprint 403 solved by the Scrapling fetcher; captures the newsroom listing, which yields little and no dates.
S&P Global — Latest News spglobal.com	sitemap	WORKING	14 / 14	2026-08-26		Listing is client-rendered and the edge refuses a headless browser (403), so neither ordinary tier can work. Read from its XML sitemap instead - working since 26 Aug (14 articles first day).
UNFCCC unfccc.int	scrape	WORKING	4 / 1	2026-08-25		Fetches and parses cleanly on schedule.
US Department of Energy energy.gov	rss	WORKING	5 / 0	2026-08-24		The energy.gov miss of 21 Aug; now healthy on RSS (5 / 7d).

The errors that occur, and what each one means

ERROR CLASS	HOW IT SHOWS	REAL EXAMPLES	REMEDY IN PLACE
TLS-fingerprint block	HTTP 403 to plain fetch and to headless Chromium alike	DCCEEW, S&P Global, EcoGeneration, NABERS	Scrapling fetcher impersonates Chrome's TLS handshake; all four now return 200
Client-rendered listing	HTTP 200 but the HTML contains no article links	S&P Global sections, vic-premier, aer-registers	S&P read from its XML sitemaps; the two XHR registers remain open items and fail visibly
Account wall	401/403 that no fetcher can pass — a genuine login	Reuters, IEA, The Australian, AFR article bodies	Newsletters subscribed into the intake mailbox (tier 0); CEC member portal has a tier-3 login
Rate limiting	HTTP 429 when a host is hit in parallel — looks identical to "no data" if unhandled	Renewables Now (23 of 31 "dateless" articles were really 429s); S&P sitemaps fetched as a burst	Per-host serialisation with delay; sitemap fetches staggered; 429 recorded as blocked and retried, never written off
No publication date in markup	Article captured but dated by fetch time — lands in the wrong day's report	7 sources in 10; measured: one article in three mis-filed	Article pages opened and article:published_time / JSON-LD read; three quality passes daily; the 23:00 pass closes the day
Unrecognised markup	Fetch reports "ok" while the parser finds nothing, for weeks	PCA (527 straight), Energy Rating/GEMS (71), CEC ("Find out more" links)	Parser fixes plus per-tier attempt logging, so "ok but empty" is visible per run
Navigation captured as news	Site furniture enters the corpus and scores well	AEMO podcast page, "IT change & release management", DWGM/ISP explainers, "Legal notice"	Model judge (news vs reference) holds standing pages out of the sheet; undecided ships
Archive unlock flood	A parser fix suddenly makes years-old pages look like today's news	40 CEC articles from 2022 (23 Aug); 484 stale-dated found on 26 Aug	report_eligible hold, applied automatically the moment a resolved date proves an article old
Scoring outage	LLM quota exhausted; unscored articles would silently vanish from the sheet	Gemini daily quota, 25 Aug	Score coalesced to 0 (ranked last, not dropped); backlog scored when quota returns
Send failure	Report renders but the mail does not arrive	DKIM failure silently discarded by Gmail (to 24 Aug); Graph 400 on semicolon recipients (26 Aug)	Send moved to Microsoft Graph (Microsoft-signed); non-202 throws so a failed send can never be recorded as sent

The two missing-article reports, diagnosed

21 August — the first complaint (~20 URLs)

Root causes, each since fixed: unrecognised markup (DCCEEW chrome filter, CEC link text, PCA and GEMS parsers — sources reported "ok" while returning nothing), TLS blocks (EcoGeneration and others, solved by the Scrapling fetcher), account walls (IEA, energy.gov then unregistered, AFR bodies — newsletter routes and new RSS sources), and listing roll-off on high-volume sites between hourly fetches (RenewEconomy, PV Magazine, Fifth Estate, Renewables Now — all now on RSS/newsletter and healthy: 51, 24, 9 and 167 articles this week respectively).

25 August — the second complaint (26 URLs)

Checked one by one against the database: 21 of 26 had been fetched all along. They were invisible for different reasons: 11 dated by fetch time into the wrong day (the date-drift fault, one in three articles), 3 cut by one point at the then-floor of 16, 2 scored 0 by an over-tight rule (since softened), 1 a physics explainer, correctly absent. The 4 S&P stories were genuinely never fetched — we monitored a different S&P section, and the right one is double-blocked (client-rendered + edge 403). Its sitemaps are now registered and delivering; the exact flagged stories were captured on the first working run.

What is read before each daily report

Hourly, every source is fetched down a five-tier ladder (email → direct → browser → login → model) with every attempt logged. Every 30 minutes new articles are scored (the scorer exits instantly when there is nothing new). Three times a day — 06:00, 12:00 and 23:00 — the quality pass opens undated articles to read their true publication date, judges news from standing pages, and writes a coverage line: fetched, dated, undated, wrong-day-corrected, and sources that ran clean but returned nothing. The 23:00 pass completes the day before it closes. At 05:00 the report queries the finished prior day — floor 1, no reference pages, no stale-dated items, minus everything already sent (the ledger guarantees nothing is lost to midnight and nothing repeats). Current volume: ~90–190 items per morning across the three brief categories.

How to capture more, and the two libraries asked about

github.com/D4Vinci/Scrapling — already integrated

This is not a proposal; it is running. The bext-scrapling container on the VPS is built on exactly this library — it is what impersonates Chrome's TLS fingerprint and turned the DCCEEW, S&P, EcoGeneration and NABERS 403s into 200s, and it fetches every article page the date-reader opens. It is tier 1 of the retrieval ladder.

github.com/firecrawl/firecrawl — worth a pilot, for a narrow job

Firecrawl self-hosts as a crawl-and-extract service with real JavaScript rendering and LLM-assisted extraction. Honest assessment: most of what it offers, the pipeline now already does — fetching, link discovery, per-page metadata, model fallback. Its genuine advantage is executing a page's JavaScript and waiting for XHR content, which is precisely the one unsolved class: vic-premier and aer-registers, and proving whether S&P natural-gas is quiet or wrong. Recommendation: pilot it self-hosted against those two or three sources as an additional ladder tier, rather than replacing a working stack wholesale. Cost: one more container (~1–2 GB RAM) on a VPS that also runs an 8B model — check headroom before committing. If the pilot solves the XHR pair, keep it for that tier alone.

Beyond the two libraries — the levers that add the most articles next

- Land the newsletters. Reuters, IEA and The Australian content only arrives by mail; the intake workflow exists — finishing the subscriptions is the cheapest coverage gain left.
- Capture the two XHR endpoints (vic-premier, aer-registers) — by hand in browser dev-tools, or by the Firecrawl pilot.
- Resolve S&P natural-gas: confirm quiet vs wrong sitemap, and register further S&P sections (electric-power, LNG) the same proven way.
- Wire the quiet-source alert to Teams. The coverage line already counts sources that ran clean and returned nothing; nobody is paged yet. This turns the next silent failure into a same-day fix instead of a client complaint.
- CEC member portal (tier 3) is built and awaiting the credential — unlocks member-only content, licensing permitting.

BEXT Consultancy · automation stack · generated 26 August 2026 · sources of truth: [sources/registry.yaml](#), live database, [fetch_attempts](#) log